

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

## Category Archives: AWS Free Dumps Online Test

26  
SEP

### Sample Questions for DAS-C01 AWS Certified Data Analytics-Specialty

To pass DAS-C01 exam and get related certification, it validates your expertise in AWS data lakes and analytics services. Build credibility and confidence by highlighting your ability to design, build, secure, and maintain analytics solutions on AWS that are efficient, cost-effective, and secure.

AWS Certified Data Analytics-Specialty DAS-C01 exam format is Multiple choice, multiple answer. You have 180 minutes to complete all the DAS-C01 questions. The available languages are English, Japanese, Korean, and Simplified Chinese.

We provide free questions of AWS Certified Data Analytics-Specialty DAS-C01 exam dumps questions, which are part of the full version. Test AWS certification DAS-C01 exam free sample questions below.

**Congratulations - you have completed this exam.**

Your answers are shown below:

**1. A smart home automation company must efficiently ingest and process messages from various connected devices and sensors. The majority of**

#### Recent Posts

Sample Questions for DAS-C01 AWS  
Certified Data Analytics-Specialty

VMware 5V0-21.19 Exam Updated Dumps  
Questions

H13-311\_v3.0 Huawei Certified ICT  
Associate-AI Dumps

2020 Updated HCIP-Access H35-211 Exam  
Dumps

[Updated in September] Certified Ethical  
Hacker 312-50v10 Dumps

# CertQueen Free Exam Dumps to Test Online


Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)


## Amazon EMR running scheduled PySpark jobs.

The data platform team manages data processing and is concerned about the efficiency and cost of downstream data processing. They want to continue to use PySpark.

**Which solution improves the efficiency of the data processing jobs and is well architected?**

Send the sensor and devices data directly to a Kinesis Data Firehose delivery stream to send the data to Amazon S3 with Apache Parquet record format conversion enabled. Use Amazon EMR running PySpark to process the data in Amazon S3. 

Set up an AWS Lambda function with a Python runtime environment. Process individual Kinesis data stream messages from the connected devices and sensors using Lambda.

**Launch an Amazon Redshift cluster. Copy the collected data from Amazon S3 to Amazon Redshift and move the data processing jobs from Amazon EMR to Amazon Redshift.** 

Set up AWS Glue Python jobs to merge the small data files in Amazon S3 into larger files and transform them to Apache Parquet format. Migrate the downstream PySpark jobs from Amazon EMR to AWS Glue.

**2. A large financial company is running its ETL process. Part of this process is to move data from Amazon S3 into an Amazon Redshift cluster. The company wants to use the most cost-efficient method to load the dataset into Amazon Redshift.**

**Which combination of steps would meet these requirements? (Choose two.)**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Use the UNLOAD command to upload data into Amazon Redshift.  
Use Amazon Redshift Spectrum to query files from Amazon S3. ✓

Explanation:

Reference: <https://aws.amazon.com/blogs/big-data/top-8-best-practices-for-high-performance-etlprocessing-using-amazon-redshift/>

**3. A university intends to use Amazon Kinesis Data Firehose to collect JSON-formatted batches of water quality readings in Amazon S3. The readings are from 50 sensors scattered across a local lake. Students will query the stored data using Amazon Athena to observe changes in a captured metric over time, such as water temperature or acidity. Interest has grown in the study, prompting the university to reconsider how data will be stored.**

**Which data format and partitioning choices will MOST significantly reduce costs? (Choose two.)**

Store the data in Apache Avro format using Snappy compression. ✓

Partition the data by year, month, and day.

Store the data in Apache ORC format using no compression. ✓

Store the data in Apache Parquet format using Snappy compression. ✓

Partition the data by sensor, year, month, and day. ✓

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/firehose/latest/dev/record-format-conversion.html>

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions


[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)


raw ERK data is stored in Amazon S3 in JSON format partitioned by hour, day, and year and is updated every hour. The company wants to maintain the data catalog and metadata in an AWS Glue Data Catalog to be able to access the data using Amazon Athena or Amazon Redshift Spectrum for analytics.

When defining tables in the Data Catalog, the company has the following requirements:

- Choose the catalog table name and do not rely on the catalog table naming algorithm.
- Keep the table updated with new partitions loaded in the respective S3 bucket prefixes.

Which solution meets these requirements with minimal effort?

Run an AWS Glue crawler that connects to one or more data stores, determines the data structures, and writes tables in the Data Catalog. 

**Use the AWS Glue console to manually create a table in the Data Catalog and schedule an AWS Lambda function to update the table partitions hourly.** 

**Use the AWS Glue API CreateTable operation to create a table in the Data Catalog. Create an AWS Glue crawler and specify the table as the source.** 

Create an Apache Hive catalog in Amazon EMR with the table schema definition in Amazon S3, and update the table partition with a scheduled job. Migrate the Hive catalog to the Data Catalog.

Explanation:

Reference: <https://docs.aws.amazon.com/glue/latest/dg/tables-described.html>



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**access must be authenticated using Microsoft Active Directory. All data in transit and at rest must be encrypted.**

## Which solution meets these requirements?

- Amazon QuickSight Standard edition configured to perform identity federation using SAML 2.0. and the default encryption settings. 
- Amazon QuickSight Enterprise edition configured to perform identity federation using SAML 2.0 and the default encryption settings.
- Amazon QuickSight Standard edition using AD Connector to authenticate using Active Directory. Configure Amazon QuickSight to use customer-provided keys imported into AWS KM
- Amazon QuickSight Enterprise edition using AD Connector to authenticate using Active Directory. Configure Amazon QuickSight to use customer-provided keys imported into AWS KM 

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/quicksight/latest/user/WhatsNew.html>

**6. An airline has been collecting metrics on flight activities for analytics. A recently completed proof of concept demonstrates how the company provides insights to data analysts to improve on-time departures. The proof of concept used objects in Amazon S3, which contained the metrics in .csv format, and used Amazon Athena for querying the data. As the amount of data increases, the data analyst wants to optimize the storage solution to improve query performance.**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

... configuration parameters.

Use an S3 bucket in the same account as Athena.

Compress the objects to reduce the data transfer I/ 

Use an S3 bucket in the same Region as Athena. 

Preprocess the .csv data to JSON to reduce I/O by fetching only the document keys needed by the query.


Preprocess the .csv data to Apache Parquet to reduce I/O by fetching only the data blocks needed for predicates.

Question was not answered

**7. A company uses the Amazon Kinesis SDK to write data to Kinesis Data Streams. Compliance requirements state that the data must be encrypted at rest using a key that can be rotated. The company wants to meet this encryption requirement with minimal coding effort.**

**How can these requirements be met?**

Create a customer master key (CMK) in AWS KM 

Assign the CMK an alias. Use the AWS Encryption SDK, providing it with the key  alias to encrypt and decrypt the data.

Create a customer master key (CMK) in AWS KM

Assign the CMK an alias. Enable server-side encryption on the Kinesis data stream using the CMK alias as the KMS master key.

Create a customer master key (CMK) in AWS KM 

Create an AWS Lambda function to encrypt and decrypt the data. Set the KMS key ID in the function's environment variables.

Enable server-side encryption on the Kinesis data stream using the default KMS key for Kinesis Data Streams.

Question was not answered

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**8. A company wants to enrich application logs in near-real-time and use the enriched dataset for further analysis. The application is running on Amazon EC2 instances across multiple Availability Zones and storing its logs using Amazon CloudWatch Logs. The enrichment source is stored in an Amazon DynamoDB table.**

**Which solution meets the requirements for the event collection and enrichment?**

Use a CloudWatch Logs subscription to send the data to Amazon Kinesis Data Firehose. Use AWS Lambda to transform the data in the Kinesis Data Firehose delivery stream and enrich it with the data in the DynamoDB table. Configure Amazon S3 as the Kinesis Data Firehose delivery destination.

Export the raw logs to Amazon S3 on an hourly basis using the AWS CL

Use AWS Glue crawlers to catalog the logs. Set up an AWS Glue connection for the DynamoDB table and set up an AWS Glue ETL job to enrich the data. Store the enriched data in Amazon S3. ✓

Configure the application to write the logs locally and use Amazon Kinesis Agent to send the data to Amazon Kinesis Data Streams. Configure a Kinesis Data Analytics SQL application with the Kinesis data stream as the source. Join the SQL application input stream with DynamoDB records, and then store the enriched output stream in Amazon S3 using Amazon Kinesis Data Firehose.

Export the raw logs to Amazon S3 on an hourly basis using the AWS CL

Use Apache Spark SQL on Amazon EMR to read the logs from Amazon S3 and enrich the records with the data from DynamoD

Store the enriched data in Amazon S3.

Question was not answered

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**A data analyst needs to load 100 billion rows of data into the new table. Before doing so, the data analyst must ensure that only members of the auditing group can read the columns containing sensitive data.**

**How can the data analyst meet these requirements with the lowest maintenance overhead?**

Load all the data into the new table and grant the auditing group permission to read from the table. Load all the data except for the columns containing sensitive data into a second table. Grant the appropriate users read-only permissions to the second table.

Load all the data into the new table and grant the auditing group permission to read from the table. Use the GRANT SQL command to allow read-only access to a subset of columns to the appropriate users.

Load all the data into the new table and grant all users read-only permissions to non-sensitive columns. Attach an IAM policy to the auditing group with explicit ALLOW access to the sensitive data columns.

Load all the data into the new table and grant the auditing group permission to read from the table. Create a view of the new table that contains all the columns, except for those considered sensitive, and grant the appropriate users read-only permissions to the table. ✓

Question was not answered

**10. A banking company wants to collect large volumes of transactional data using Amazon Kinesis Data Streams for real-time analytics. The company uses PutRecord to send data to Amazon Kinesis, and has observed network outages during certain times of the day. The company wants to obtain exactly once semantics for the entire processing pipeline.**



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Processing a single record in each stream.

Rely on the processing semantics of Amazon Kinesis Data Analytics to avoid duplicate processing of events.

Design the data producer so events are not ingested into Kinesis Data Streams multiple times.

Rely on the exactly one processing semantics of Apache Flink and Apache Spark Streaming included in Amazon EM

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/streams/latest/dev/kinesis-record-processor-duplicates.html>

**11. A company's data analyst needs to ensure that queries executed in Amazon Athena cannot scan more than a prescribed amount of data for cost control purposes. Queries that exceed the prescribed threshold must be canceled immediately.**

**What should the data analyst do to achieve this?**

Configure Athena to invoke an AWS Lambda function that terminates queries when the prescribed threshold is crossed. ✓

For each workgroup, set the control limit for each query to the prescribed threshold.

Enforce the prescribed threshold on all Amazon S3 bucket policies

For each workgroup, set the workgroup-wide data usage control limit to the prescribed threshold. ✓

Question was not answered


# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**12. A marketing company is using Amazon EMR clusters for its workloads. The company manually installs third-party libraries on the clusters by logging in to the master nodes. A data analyst needs to create an automated solution to replace the manual process.**

**Which options can fulfill these requirements? (Choose two.)**

Place the required installation scripts in Amazon S3 and execute them using custom bootstrap actions. 

Place the required installation scripts in Amazon S3 and execute them through Apache Spark in Amazon EM

Install the required third-party libraries in the existing EMR master node. Create an AMI out of that master node and use that custom AMI to re-create the EMR cluster.

Use an Amazon DynamoDB table to store the list of required applications. Trigger an AWS Lambda function with DynamoDB Streams to install the software.

Launch an Amazon EC2 instance with Amazon Linux and install the required third-party libraries on the instance. Create an AMI and use that AMI to create the EMR cluster.

Question was not answered

**13. A data engineering team within a shared workspace company wants to build a centralized logging system for all weblogs generated by the space reservation system. The company has a fleet of Amazon EC2 instances that process requests for shared space reservations on its website. The data engineering team wants to ingest all weblogs into a service that will provide a near-real-time search engine. The team does not want to manage the**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Set up the Amazon CloudWatch agent to stream weblogs to CloudWatch logs and subscribe the Amazon Kinesis data stream to CloudWatch. Choose Amazon Elasticsearch Service as the end destination of the weblogs. Set up the Amazon CloudWatch agent to stream weblogs to CloudWatch logs and subscribe the Amazon Kinesis Data Firehose delivery stream to CloudWatch. Choose Amazon Elasticsearch Service as the end destination of the weblogs. Set up the Amazon CloudWatch agent to stream weblogs to CloudWatch logs and subscribe the Amazon Kinesis data stream to CloudWatch. Configure Splunk as the end destination of the weblogs. Set up the Amazon CloudWatch agent to stream weblogs to CloudWatch logs and subscribe the Amazon Kinesis Firehose delivery stream to CloudWatch. Configure Amazon DynamoDB as the end destination of the weblogs.

Question was not answered

**14. A company wants to research user turnover by analyzing the past 3 months of user activities. With millions of users, 1.5 TB of uncompressed data is generated each day. A 30-node Amazon Redshift cluster with 2.56 TB of solid state drive (SSD) storage for each node is required to meet the query performance goals.**

**The company wants to run an additional analysis on a year's worth of historical data to examine trends indicating which features are most popular. This analysis will be done once a week.**

**What is the MOST cost-effective solution?**

Increase the size of the Amazon Redshift cluster to 120 nodes so it has enough storage capacity to hold 1 year of data. Then use Amazon Redshift for the additional analysis.

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

days to Amazon S3 and store it in Apache Parquet format partitioned by date. Then provision a persistent Amazon EMR cluster and use Apache Presto for the additional analysis.

Resize the cluster node type to the dense storage node type (DS2) for an additional 16 TB storage capacity on each individual node in the Amazon Redshift cluster. Then use Amazon Redshift for the additional analysis.

Question was not answered

**15. A bank operates in a regulated environment. The compliance requirements for the country in which the bank operates say that customer data for each state should only be accessible by the bank's employees located in the same state. Bank employees in one state should NOT be able to access data for customers who have provided a home address in a different state.**

**The bank's marketing team has hired a data analyst to gather insights from customer data for a new campaign being launched in certain states. Currently, data linking each customer account to its home state is stored in a tabular .csv file within a single Amazon S3 folder in a private S3 bucket. The total size of the S3 folder is 2 GB uncompressed. Due to the country's compliance requirements, the marketing team is not able to access this folder.**

**The data analyst is responsible for ensuring that the marketing team gets one-time access to customer data for their campaign analytics project, while being subject to all the compliance requirements and controls.**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions


[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

marketing employees with appropriate data access under compliance controls.

Delete the bucket policies after the project.

Load tabular data from Amazon S3 to an Amazon EMR cluster using s3DistCp.

Implement a custom Hadoop-based row-level security solution on the Hadoop Distributed File System (HDFS) to provide marketing employees with appropriate data access under compliance controls. Terminate the EMR cluster after the project.

Load tabular data from Amazon S3 to Amazon Redshift with the COPY command. Use the built-in row-level security feature in Amazon Redshift to provide marketing employees with appropriate data access under compliance controls. Delete the Amazon Redshift tables after the project. 

Load tabular data from Amazon S3 to Amazon QuickSight Enterprise edition by directly importing it as a data source. Use the built-in row-level security feature in Amazon QuickSight to provide marketing employees with appropriate data access under compliance controls. Delete Amazon QuickSight data sources after the project is complete.

Question was not answered

**16. A financial services company needs to aggregate daily stock trade data from the exchanges into a data store. The company requires that data be streamed directly into the data store, but also occasionally allows data to be modified using SQL. The solution should integrate complex, analytic queries running with minimal latency. The solution must provide a business intelligence dashboard that enables viewing of the top contributors to anomalies in stock prices.**

**Which solution meets the company's requirements?**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Amazon Redshift as a data source for Amazon QuickSight to create a business intelligence dashboard.

Use Amazon Kinesis Data Firehose to stream data to Amazon Redshift. Use Amazon Redshift as a data source for Amazon QuickSight to create a business intelligence dashboard.

Use Amazon Kinesis Data Streams to stream data to Amazon S3. Use Amazon Athena as a data source for Amazon QuickSight to create a business intelligence dashboard. ✓

Question was not answered

**17. A financial company hosts a data lake in Amazon S3 and a data warehouse on an Amazon Redshift cluster. The company uses Amazon QuickSight to build dashboards and wants to secure access from its on-premises Active Directory to Amazon QuickSight.**

**How should the data be secured?**

Use an Active Directory connector and single sign-on (SSO) in a corporate network environment.

Use a VPC endpoint to connect to Amazon S3 from Amazon QuickSight and an IAM role to authenticate Amazon Redshift. ✓

Establish a secure connection by creating an S3 endpoint to connect Amazon QuickSight and a VPC endpoint to connect to Amazon Redshift.

Place Amazon QuickSight and Amazon Redshift in the security group and use an Amazon S3 endpoint to connect Amazon QuickSight to Amazon S3.

Question was not answered


# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**File System (HDFS). The company wants a cost-effective solution to make its HBase data highly available.**

## Which architectural pattern meets company's requirements?

Use Spot Instances for core and task nodes and a Reserved Instance for the EMR master node. Configure the EMR cluster with multiple master nodes. 


Schedule automated snapshots using Amazon EventBridge.

Store the data on an EMR File System (EMRFS) instead of HDF

Enable EMRFS consistent view. 

Create an EMR HBase cluster with multiple master nodes. Point the HBase root directory to an Amazon S3 bucket.

Store the data on an EMR File System (EMRFS) instead of HDFS and enable EMRFS consistent view. Run two separate EMR clusters in two different Availability Zones. Point both clusters to the same HBase root directory in the same Amazon S3 bucket.

Store the data on an EMR File System (EMRFS) instead of HDFS and enable EMRFS consistent view. Create a primary EMR HBase cluster with multiple master nodes. Create a secondary EMR HBase read-replica cluster in a separate Availability Zone. Point both clusters to the same HBase root directory in the same Amazon S3 bucket. 

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/emr/latest/ReleaseGuide/emr-hbase-s3.html>

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**that the application is working as expected after each deployment. The collection and analysis solution should be highly available with the ability to display new information with minimal delays.**

**Which method should the company use to collect and analyze the logs?**

Enable detailed monitoring on Amazon EC2, use Amazon CloudWatch agent to store logs in Amazon S3, and use Amazon Athena for fast, interactive log analytics. ✓

Use the Amazon Kinesis Producer Library (KPL) agent on Amazon EC2 to collect and send data to Kinesis Data Streams to further push the data to Amazon Elasticsearch Service and visualize using Amazon QuickSight.

Use the Amazon Kinesis Producer Library (KPL) agent on Amazon EC2 to collect and send data to Kinesis Data Firehose to further push the data to Amazon Elasticsearch Service and Kibana. ✓

Use Amazon CloudWatch subscriptions to get access to a real-time feed of logs and have the logs delivered to Amazon Kinesis Data Streams to further push the data to Amazon Elasticsearch Service and Kibana. ✓

Question was not answered

Explanation:

Reference:

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/Subscriptions.html>

**20. A data analyst is using AWS Glue to organize, cleanse, validate, and format a 200 GB dataset. The data analyst triggered the job to run with the Standard worker type. After 3 hours, the AWS Glue job status is still RUNNING. Logs from the job run show no error codes. The data analyst**



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Enable job bookmarks in AWS Glue to estimate the number of data processing units (DPUs). Based on the profiled metrics, increase the value of the executor-cores job parameter. ✓

Enable job metrics in AWS Glue to estimate the number of data processing units (DPUs). Based on the profiled metrics, increase the value of the maximum capacity job parameter. ✓

Enable job metrics in AWS Glue to estimate the number of data processing units (DPUs). Based on the profiled metrics, increase the value of the spark.yarn.executor.memoryOverhead job parameter.

Enable job bookmarks in AWS Glue to estimate the number of data processing units (DPUs). Based on the profiled metrics, increase the value of the num-executors job parameter.

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/glue/latest/dg/monitor-debug-capacity.html>

**21. A company has a business unit uploading .csv files to an Amazon S3 bucket. The company's data platform team has set up an AWS Glue crawler to do discovery, and create tables and schemas. An AWS Glue job writes processed data from the created tables to an Amazon Redshift database. The AWS Glue job handles column mapping and creating the Amazon Redshift table appropriately. When the AWS Glue job is rerun for any reason in a day, duplicate records are introduced into the Amazon Redshift table.**

**Which solution will update the Redshift table without duplicates when jobs are rerun?**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Perform an upsert operation in MySQL, and copy the results to the Amazon Redshift table.

Use Apache Spark's DataFrame dropDuplicates() API to eliminate duplicates and then write the data to Amazon Redshift.

Use the AWS Glue ResolveChoice built-in transform to select the most recent value of the column.

Question was not answered

Explanation:

Reference: <https://towardsdatascience.com/update-and-insert-upsert-data-from-aws-glue-698ac582e562>

**22. A streaming application is reading data from Amazon Kinesis Data Streams and immediately writing the data to an Amazon S3 bucket every 10 seconds. The application is reading data from hundreds of shards. The batch interval cannot be changed due to a separate requirement. The data is being accessed by Amazon Athena. Users are seeing degradation in query performance as time progresses.**

**Which action can help improve query performance?**

Merge the files in Amazon S3 to form larger files.

Increase the number of shards in Kinesis Data Streams.

Add more memory and CPU capacity to the streaming application. 

Write the files to multiple S3 buckets.

Question was not answered

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

an Amazon ES cluster.

The company has very slow query performance on the Amazon ES index and occasionally sees errors from Kinesis Data Firehose when attempting to write to the index. The Amazon ES cluster has 10 nodes running a single index and 3 dedicated master nodes. Each data node has 1.5 TB of Amazon EBS storage attached and the cluster is configured with 1,000 shards. Occasionally, JVMMemoryPressure errors are found in the cluster logs.

**Which solution will improve the performance of Amazon ES?**

Increase the memory of the Amazon ES master nodes.

Decrease the number of Amazon ES data nodes.

Decrease the number of Amazon ES shards for the index. ✓

Increase the number of Amazon ES shards for the index.

Question was not answered

**24. A manufacturing company has been collecting IoT sensor data from devices on its factory floor for a year and is storing the data in Amazon Redshift for daily analysis. A data analyst has determined that, at an expected ingestion rate of about 2 TB per day, the cluster will be undersized in less than 4 months. A long-term solution is needed. The data analyst has indicated that most queries only reference the most recent 13 months of data, yet there are also quarterly reports that need to query all the data generated from the past 7 years. The chief technology officer (CTO) is concerned about the costs, administrative effort, and performance of a long-term solution.**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

table in Amazon Redshift to point to the S3 location. Use Amazon Redshift Spectrum to join to data that is older than 13 months.

Take a snapshot of the Amazon Redshift cluster. Restore the cluster to a new cluster using dense storage nodes with additional storage capacity. ✓

Execute a CREATE TABLE AS SELECT (CTAS) statement to move records that are older than 13 months to quarterly partitioned data in Amazon Redshift Spectrum backed by Amazon S3.

Unload all the tables in Amazon Redshift to an Amazon S3 bucket using S3 Intelligent-Tiering. Use AWS Glue to crawl the S3 bucket location to create external tables in an AWS Glue Data Catalog. Create an Amazon EMR cluster using Auto Scaling for any daily analytics needs, and use Amazon Athena for the quarterly reports, with both using the same AWS Glue Data Catalog.

Question was not answered

**25. An insurance company has raw data in JSON format that is sent without a predefined schedule through an Amazon Kinesis Data Firehose delivery stream to an Amazon S3 bucket. An AWS Glue crawler is scheduled to run every 8 hours to update the schema in the data catalog of the tables stored in the S3 bucket. Data analysts analyze the data using Apache Spark SQL on Amazon EMR set up with AWS Glue Data Catalog as the metastore. Data analysts say that, occasionally, the data they receive is stale. A data engineer needs to provide access to the most up-to-date data.**

**Which solution meets these requirements?**

Create an external schema based on the AWS Glue Data Catalog on the existing Amazon Redshift cluster to query new data in Amazon S3 with Amazon Redshift Spectrum. ✓

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Run the AWS Glue crawler from an AWS Lambda function triggered by an S3:ObjectCreated:\* event notification on the S3 bucket.


Question was not answered

**26. A company that produces network devices has millions of users. Data is collected from the devices on an hourly basis and stored in an Amazon S3 data lake.**

**The company runs analyses on the last 24 hours of data flow logs for abnormality detection and to troubleshoot and resolve user issues. The company also analyzes historical logs dating back 2 years to discover patterns and look for improvement opportunities.**

**The data flow logs contain many metrics, such as date, timestamp, source IP, and target IP. There are about 10 billion events every day.**

**How should this data be stored for optimal performance?**

- In Apache ORC partitioned by date and sorted by source IP
- In compressed .csv partitioned by date and sorted by source IP
- In Apache Parquet partitioned by source IP and sorted by date
- In compressed nested JSON partitioned by source IP and sorted by date 

Question was not answered

**27. A banking company is currently using an Amazon Redshift cluster with dense storage (DS) nodes to store sensitive data. An audit found that the**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

**Which combination of steps is required to achieve compliance? (Choose two.)**

Set up a trusted connection with HSM using a client and server certificate with automatic key rotation. ✓

Modify the cluster with an HSM encryption option and automatic key rotation. ✓

Create a new HSM-encrypted Amazon Redshift cluster and migrate the data to the new cluster.

Enable HSM with key rotation through the AWS CL ✓

Enable Elliptic Curve Diffie-Hellman Ephemeral (ECDHE) encryption in the HS

Question was not answered

Explanation:

Reference: <https://docs.aws.amazon.com/redshift/latest/mgmt/working-with-db-encryption.html>

**28. A company is planning to do a proof of concept for a machine learning (ML) project using Amazon SageMaker with a subset of existing on-premises data hosted in the company's 3 TB data warehouse. For part of the project, AWS Direct Connect is established and tested. To prepare the data for ML, data analysts are performing data curation. The data analysts want to perform multiple steps, including mapping, dropping null fields, resolving choice, and splitting fields. The company needs the fastest solution to curate the data for this project.**


**Which solution meets these requirements?**

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

data into Amazon S3 for ML processing.

Ingest data into Amazon S3 using AWS DM 

Use AWS Glue to perform data curation and store the data in Amazon 3 for ML processing.

Take a full backup of the data store and ship the backup files using AWS Snowball. Upload Snowball data into Amazon S3 and schedule data curation jobs using AWS Batch to prepare the data for M


Question was not answered

**29. A US-based sneaker retail company launched its global website. All the transaction data is stored in Amazon RDS and curated historic transaction data is stored in Amazon Redshift in the us-east-1 Region. The business intelligence (BI) team wants to enhance the user experience by providing a dashboard for sneaker trends.**

**The BI team decides to use Amazon QuickSight to render the website dashboards. During development, a team in Japan provisioned Amazon QuickSight in ap-northeast-1. The team is having difficulty connecting Amazon QuickSight from ap-northeast-1 to Amazon Redshift in us-east-1.**

**Which solution will solve this issue and meet the requirements?**

In the Amazon Redshift console, choose to configure cross-Region snapshots and set the destination Region as ap-northeast-1. Restore the Amazon Redshift Cluster from the snapshot and connect to Amazon QuickSight launched in ap-northeast-1.

Create a VPC endpoint from the Amazon QuickSight VPC to the Amazon Redshift VPC so Amazon QuickSight can access data from Amazon Redshift. 

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

rule authorizing access from the appropriate IP address range for the Amazon QuickSight servers in ap-northeast-1.

Question was not answered

**30. An airline has .csv-formatted data stored in Amazon S3 with an AWS Glue Data Catalog. Data analysts want to join this data with call center data stored in Amazon Redshift as part of a daily batch process. The Amazon Redshift cluster is already under a heavy load. The solution must be managed, serverless, well-functioning, and minimize the load on the existing Amazon Redshift cluster. The solution should also require minimal effort and development activity.**

**Which solution meets these requirements?**

Unload the call center data from Amazon Redshift to Amazon S3 using an AWS Lambda function. Perform the join with AWS Glue ETL scripts.

Export the call center data from Amazon Redshift using a Python shell in AWS Glue. Perform the join

with AWS Glue ETL scripts.

Create an external table using Amazon Redshift Spectrum for the call center data and perform the join with Amazon Redshift. ✓

Export the call center data from Amazon Redshift to Amazon EMR using Apache Sqoop. Perform the join with Apache Hive.

Question was not answered





# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

Leave a comment

12  
SEP

## DVA-C01 AWS Certified Developer–Associate Exam Dumps

AWS Certified Developer–Associate DVA-C01 exam dumps have been cracked, which are valuable for you to prepare the test. AWS Certified Developer–Associate DVA-C01 exam is intended for individuals who perform a development role and have one or more years of hands-on experience developing and maintaining an AWS-based application.

AWS DVA-C01 exam format is Multiple choice, multiple answer. You have 130 minutes to complete all the DVA-C01 questions. The exam languages are English, Japanese, Korean, and Simplified Chinese. DVA-C01 exam cost is 150 USD.

We provide free questions of AWS certification DVA-C01 exam dumps, which are part of the full version. Study AWS Certified Developer–Associate DVA-C01 exam free dumps below.

Please go to [DVA-C01 AWS Certified Developer–Associate Exam Dumps](#) to view this quiz



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

28  
AUG

## AWS Certified Solutions Architect– Professional SAP-C01 Exam Dumps

The AWS Certified Solutions Architect–Professional SAP-C01 exam is intended for individuals who perform a solutions architect role with two or more years of hands-on experience managing and operating systems on AWS.

We have cracked the latest AWS certification SAP-C01 exam dumps, which are valuable for you to pass SAP-C01 test. Each candidate has 180 minutes to complete the test. The exam cost is \$300. SAP-C01 test languages are English, Japanese, Korean, and Simplified Chinese.

We provide free questions of AWS Certified Solutions Architect–Professional SAP-C01 exam dumps, which are part of the full version. Practice AWS certification SAP-C01 exam free dumps below.

Please go to [AWS Certified Solutions Architect–Professional SAP-C01 Exam Dumps](#) to view this quiz



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

26  
AUG

## AWS Certified Developer-Associate DVA-C01 Exam Dumps Questions

AWS Certified Developer-Associate DVA-C01 exam is intended for individuals who perform a development role and have one or more years of hands-on experience developing and maintaining an AWS-based application.

You have 130 minutes to complete AWS Certified Developer-Associate DVA-C01 test. The exam questions format is Multiple choice, multiple answer. DVA-C01 exam languages are English, Japanese, Korean, and Simplified Chinese.

We provide free questions of AWS Certified Developer-Associate DVA-C01 exam dumps questions, which are part of the full version. Practice AWS certification DVA-C01 exam free dumps questions below.

Please go to [AWS Certified Developer-Associate DVA-C01 Exam Dumps Questions](#) to view this quiz



admin



Posted in AWS Free Dumps Online Test



AWS certification

DVA-C01 exam, AWS Certified Developer-Associate DVA-C01 exam, DVA-C01 exam

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

20  
AUG





## Passed SAA-C02 Exam with Score 775

AWS Certified Solutions Architect-Associate SAA-C02 dumps with 189 Q&As are valid. We have customers just passed SAA-C02 exam with score 775. If you are preparing your AWS certification SAA-C02 exam, the updated SAA-C02 exam dumps are the best material.

We provide free questions of AWS certification SAA-C02 exam dumps, which are part of the full version. Practice the updated AWS Certified Solutions Architect-Associate SAA-C02 exam free dumps below.

Please go to [Passed SAA-C02 Exam with Score 775](#) to view this quiz



 admin  Posted in AWS Free Dumps Online Test  AWS certification SAA-C02 exam dumps, AWS Certified Solutions Architect-Associate SAA-C02 exam, SAA-C02 exam dumps, SAA-C02 exam free dumps  Leave a comment

18  
AUG

## CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)





The AWS Certified DevOps Engineer-Professional DOP-C01 exam is intended for individuals who perform a DevOps engineer role with two or more years of experience provisioning, operating, and managing AWS environments.

You have 180 minutes to complete AWS certification DOP-C01 exam. The exam format is Multiple choice and multiple answer. DOP-C01 exam languages are English, Japanese, Korean, and Simplified Chinese.

We provide free questions of AWS Certified DevOps Engineer-Professional DOP-C01 dumps questions, which are part of the full version. Test AWS certification DOP-C01 exam free dumps questions below.

Please go to [AWS Certified DevOps Engineer-Professional DOP-C01 Dumps Questions](#) to view this quiz



 admin  Posted in AWS Free Dumps Online Test  AWS certification DOP-C01 exam, AWS Certified DevOps Engineer-Professional, DOP-C01 exam, DOP-C01 exam free dumps questions  Leave a comment

15  
AUG

## AWS Certified Advanced Networking-Specialty ANS-C00 Exam Updated Dumps

## CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)





who perform complex networking tasks. We have updated AWS certification AWS C00 exam dumps, which are valuable for you to pass ANS-C00 exam.

Candidates have 170 minutes to complete all the ANS-C00 AWS Certified Advanced Networking-Specialty exam questions. ANS-C00 exam format is Multiple choice, multiple answer. You need to spend \$300 to register the test. ANS-C00 exam languages are English, Japanese, Korean, and Simplified Chinese.

We provide free updated questions of AWS Certified Advanced Networking-Specialty ANS-C00 exam dumps, which are part of the full version. Practice AWS certified ANS-C00 exam free updated dumps below.

Please go to [AWS Certified Advanced Networking-Specialty ANS-C00 Exam Updated Dumps](#) to view this quiz



 admin  Posted in AWS Free Dumps Online Test  ANS-C00 exam, ANS-C00 exam dumps, ANS-C00 exam free updated dumps, AWS certification ANS-C00 exam dumps, AWS Certified Advanced Networking-Specialty ANS-C00, AWS certified ANS-C00 exam  Leave a comment

29  
JUL

## AWS Certified Solutions Architect-Associate SAA-C02 Dumps Questions

# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions





[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

SAA-C02 exam validates your ability to define a solution using architectural design principles based on customer requirements and provide implementation guidance based on best practices to an organization throughout the lifecycle of a project.

We provide free questions of AWS Certified Solutions Architect-Associate SAA-C02 exam dumps questions, which are part of the full version. Test AWS certification SAA-C02 exam free dumps questions below.

Please go to [AWS Certified Solutions Architect-Associate SAA-C02 Dumps Questions](#) to view this quiz



 admin  Posted in AWS Free Dumps Online Test  AWS Certified Solutions Architect-Associate SAA-C02 exam, SAA-C02 exam, SAA-C02 exam dumps questions, SAA-C02 exam free dumps questions  Leave a comment

20  
JUL

## DAS-C01 AWS Certified Data Analytics-Specialty Exam Dumps

Passing AWS Certified Data Analytics-Specialty DAS-C01 exam is not difficult now. We have cracked the latest AWS certification DAS-C01 exam dumps, which are

## CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions

[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

maintain analytics solutions that provide insight from data.

We provide free questions of AWS certification DAS-C01 exam dumps, which are part of the full version. Study the following AWS Certified DAS-C01 exam free dumps below.

Please go to [DAS-C01 AWS Certified Data Analytics-Specialty Exam Dumps](#) to view this quiz



admin Posted in AWS Free Dumps Online Test AWS certification DAS-C01 exam dumps, AWS certified DAS-C01 exam, AWS Certified Data Analytics-Specialty DAS-C01 exam, DAS-C01 exam dumps, DAS-C01 exam free dumps Leave a comment

06  
JUN

## Updated AWS Certified Solutions Architect-Professional Exam Dumps Questions

AWS Certified Solutions Architect-Professional exam validates your technical skills and expertise with an industry-recognized credential and grow your career. To



# CertQueen Free Exam Dumps to Test Online

Platform to offer free real exam questions




[HOME](#)[HUAWEI FREE DUMPS](#)[CISCO FREE DUMPS](#)

individuals who perform a solutions architect role with two or more years of hands-on experience managing and operating systems on AWS.

We provide free questions of AWS Certified Solutions Architect-Professional exam dumps questions, which are part of the full version. Practice AWS Certified Solutions Architect-Professional exam free dumps questions below.

Please go to [Updated AWS Certified Solutions Architect-Professional Exam Dumps Questions](#) to view this quiz



admin  Posted in [AWS Free Dumps Online Test](#)  [AWS Certified Solutions Architect](#), [AWS Certified Solutions Architect-Professional exam](#), [AWS Certified Solutions Architect-Professional exam dumps questions](#), [AWS Certified Solutions Architect-Professional exam free dumps questions](#)  4 Comments

[← Older posts](#)

Fruitful theme by [fruitfulcode](#) Powered by: [WordPress](#)

